

1. Does it make sense to calculate the p-value on individual regression coefficients from a ridge or lasso model? Why or why not?
2. The traditional notion of the bias-variance tradeoff states that when the number of parameters in your model, p , exceeds the number of observations, n , there is no unique solution to your model. How exactly then does ridge regression, which does not remove variables from your model, resolve the bias-variance tradeoff when $p > n$?