

Bagging, random forests, & boosting

Readings for today

- Chapter 8: Tree-based methods. James, G., Witten, D., Hastie, T., & Tibshirani, R. (2013). An introduction to statistical learning: with applications in R (Vol. 6). New York: Springer

Topics

1. Bagging

2. Random forests

3. Boosting

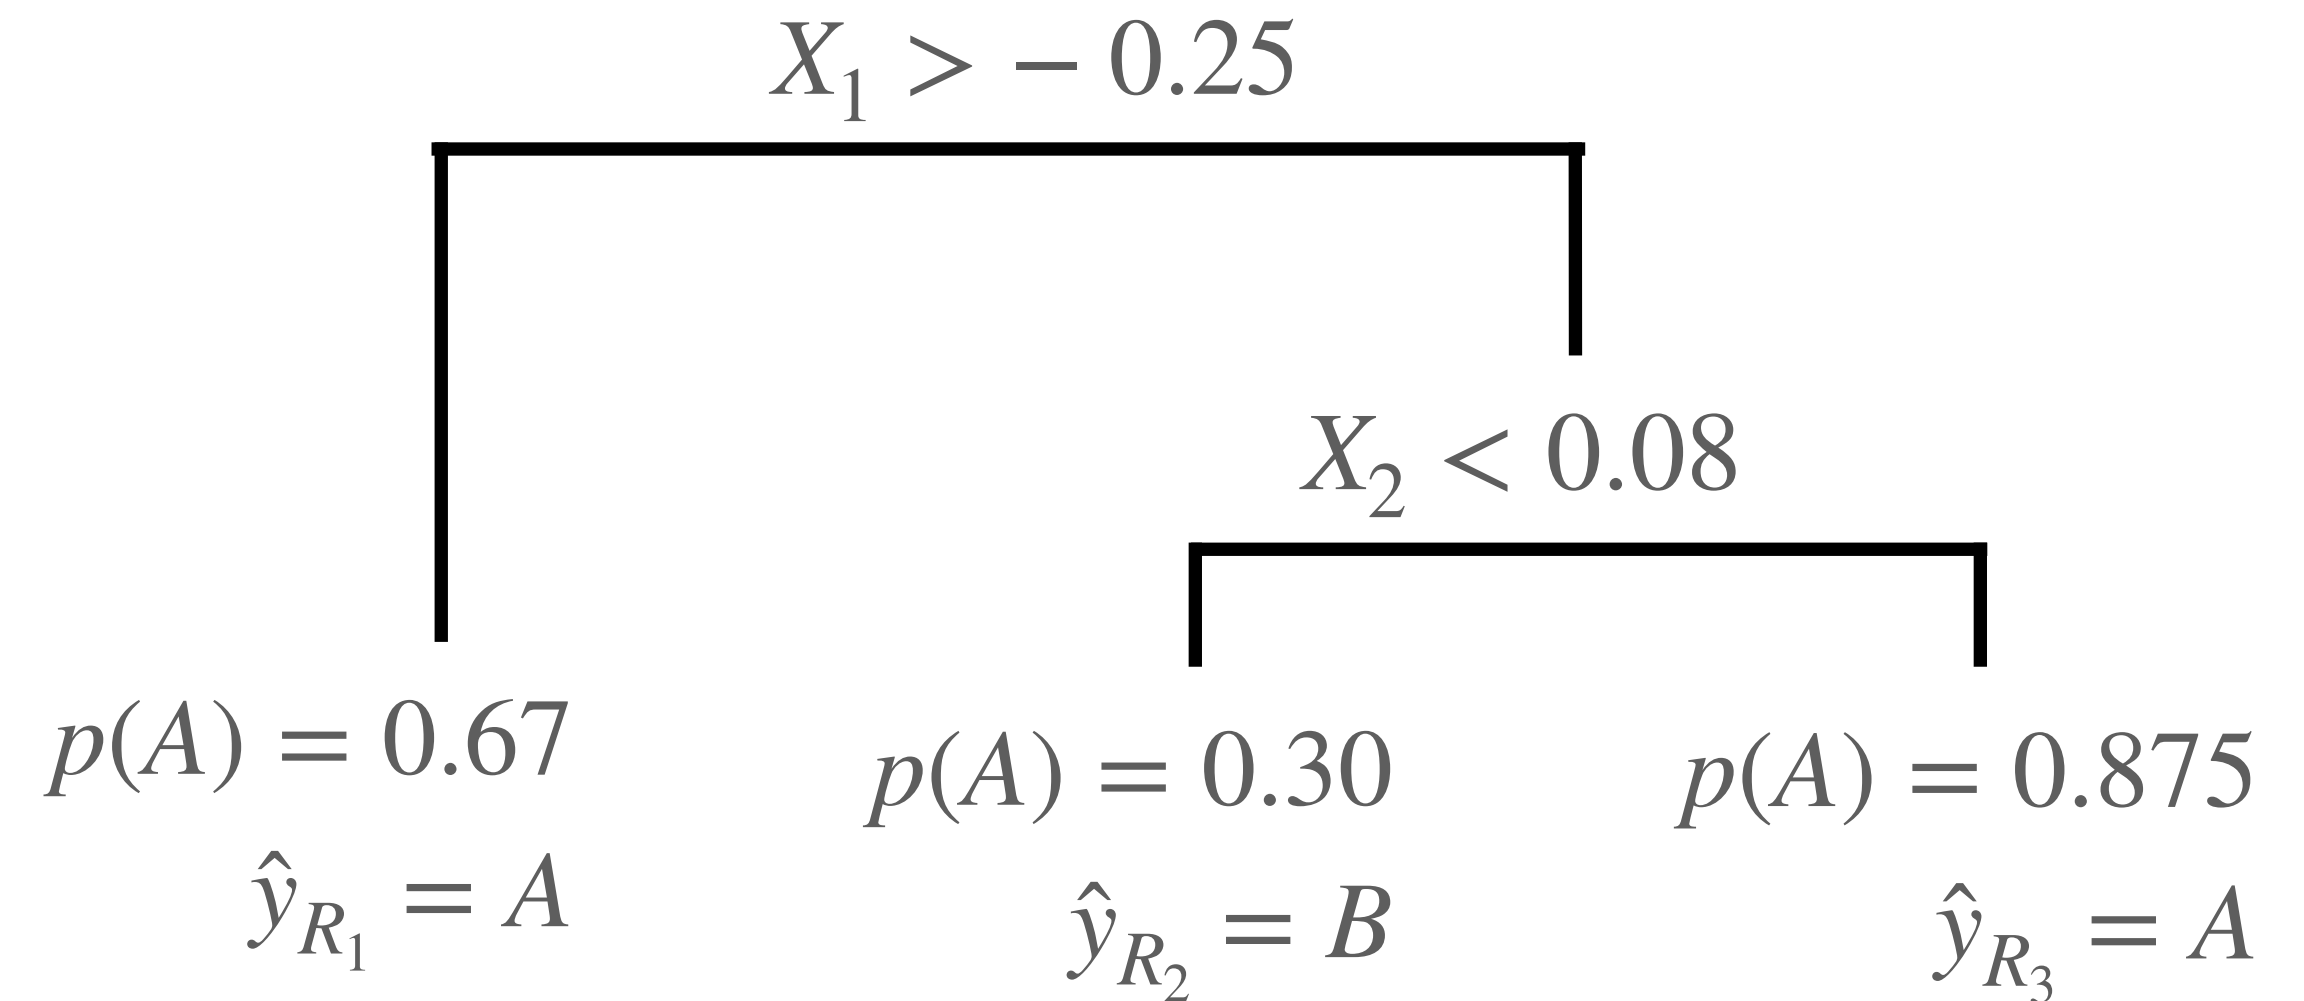
Bagging

Problems with decision trees

Disadvantages:

- Lower predictive accuracy.
- Non-robust

small changes in data have huge impacts on model fits.



High flexibility of decision trees leads to sensitivity to variation in data set.

Bootstrap aggregation (Bagging)

$$\begin{array}{c} \overbrace{\begin{pmatrix} y_1 \\ \vdots \\ y_n \end{pmatrix}}^{Y_1^*} \quad \overbrace{\begin{pmatrix} x_{1,1} & \cdots & x_{1,p} \\ \vdots & & \vdots \\ x_{n,1} & \cdots & x_{n,p} \end{pmatrix}}^{X_1^*} \rightarrow \hat{f}^1(X_1^*) \\ \vdots \\ \overbrace{\begin{pmatrix} y_1 \\ \vdots \\ y_n \end{pmatrix}}^{Y_B^*} \quad \overbrace{\begin{pmatrix} x_{1,1} & \cdots & x_{1,p} \\ \vdots & & \vdots \\ x_{n,1} & \cdots & x_{n,p} \end{pmatrix}}^{X_B^*} \rightarrow \hat{f}^B(X_B^*) \end{array} \quad \left| \quad \hat{y}_i = \frac{1}{B} \sum_{b=1}^B \hat{f}^b(x_i) \right.$$

- Use bootstrapping to generate B predictive models.
- For any prediction of y_i , run x_i through all B models.
- Use the average of those predictions for \hat{y}_i .

Example: regression

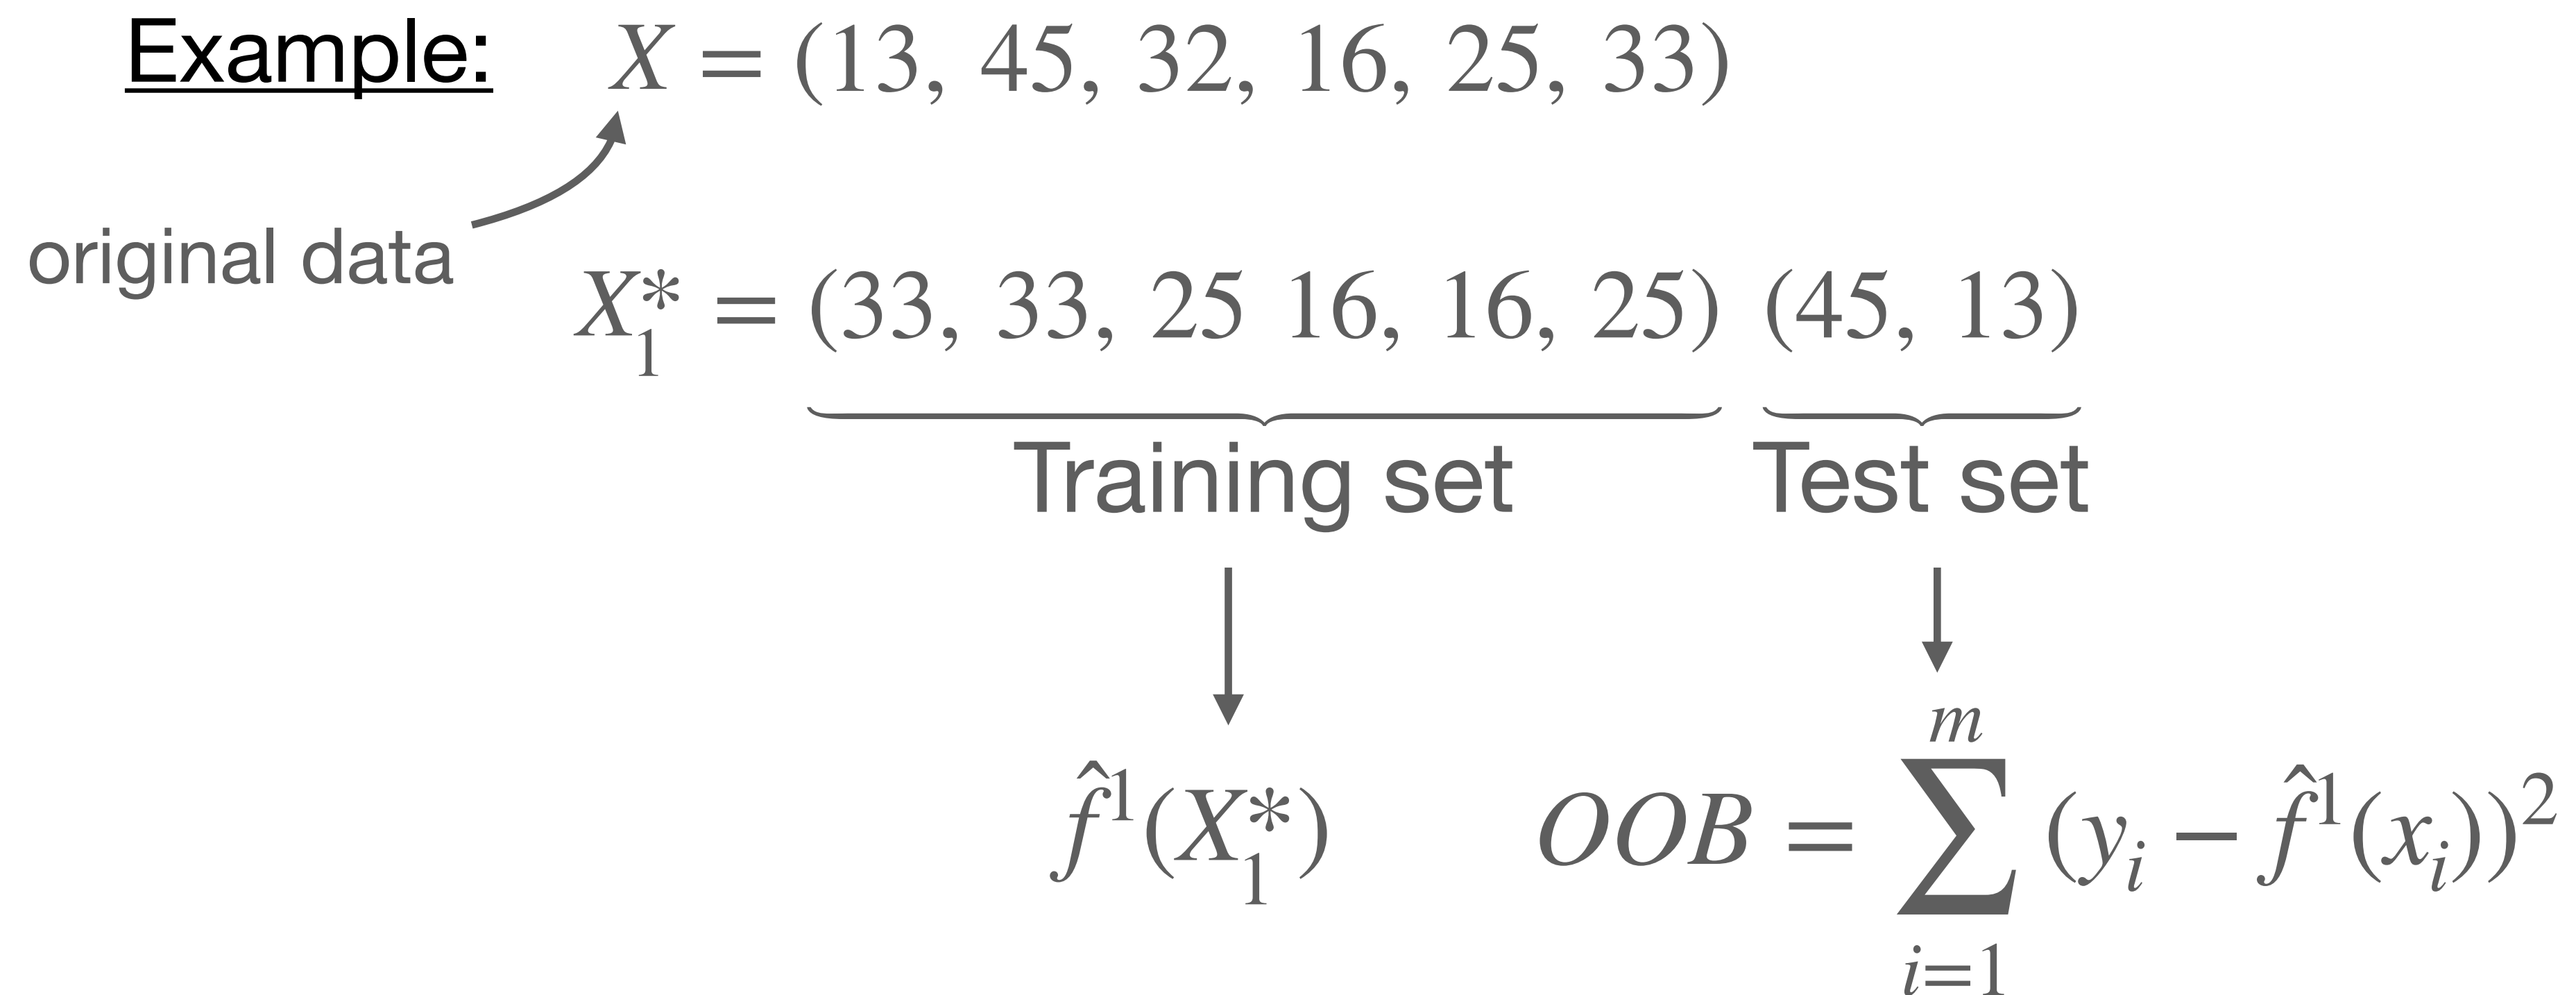
$$\begin{array}{l}
 \overbrace{\begin{pmatrix} y_1 \\ \vdots \\ y_n \end{pmatrix}}^{Y_1^*} = \overbrace{\begin{pmatrix} x_{1,1} & \cdots & x_{1,p} \\ \vdots & & \vdots \\ x_{n,1} & \cdots & x_{n,p} \end{pmatrix}}^{X_1^*} \overbrace{\begin{pmatrix} \hat{\beta}_1^* \\ \vdots \\ \hat{\beta}_p^* \end{pmatrix}}^{\hat{\beta}_1^*} \rightarrow \hat{f}^1(X_1^*) = \hat{\beta}_{1,1}^* X_{1,1}^* + \cdots + \hat{\beta}_{1,p}^* X_{1,p}^* \\
 \vdots \\
 \overbrace{\begin{pmatrix} y_1 \\ \vdots \\ y_n \end{pmatrix}}^{Y_B^*} = \overbrace{\begin{pmatrix} x_{1,1} & \cdots & x_{1,p} \\ \vdots & & \vdots \\ x_{n,1} & \cdots & x_{n,p} \end{pmatrix}}^{X_B^*} \overbrace{\begin{pmatrix} \hat{\beta}_1^* \\ \vdots \\ \hat{\beta}_p^* \end{pmatrix}}^{\hat{\beta}_B^*} \rightarrow \hat{f}^B(X_B^*) = \hat{\beta}_{B,1}^* X_{1,1}^* + \cdots + \hat{\beta}_{B,p}^* X_{B,p}^*
 \end{array}
 \left| \begin{array}{l}
 \hat{y}_i = \frac{1}{B} \sum_{b=1}^B \hat{f}^b(x_i) \\
 = \frac{1}{B} \sum_{b=1}^B \sum_{j=1}^p \hat{\beta}_{b,j}^* x_{i,j}
 \end{array} \right.$$

Example: trees

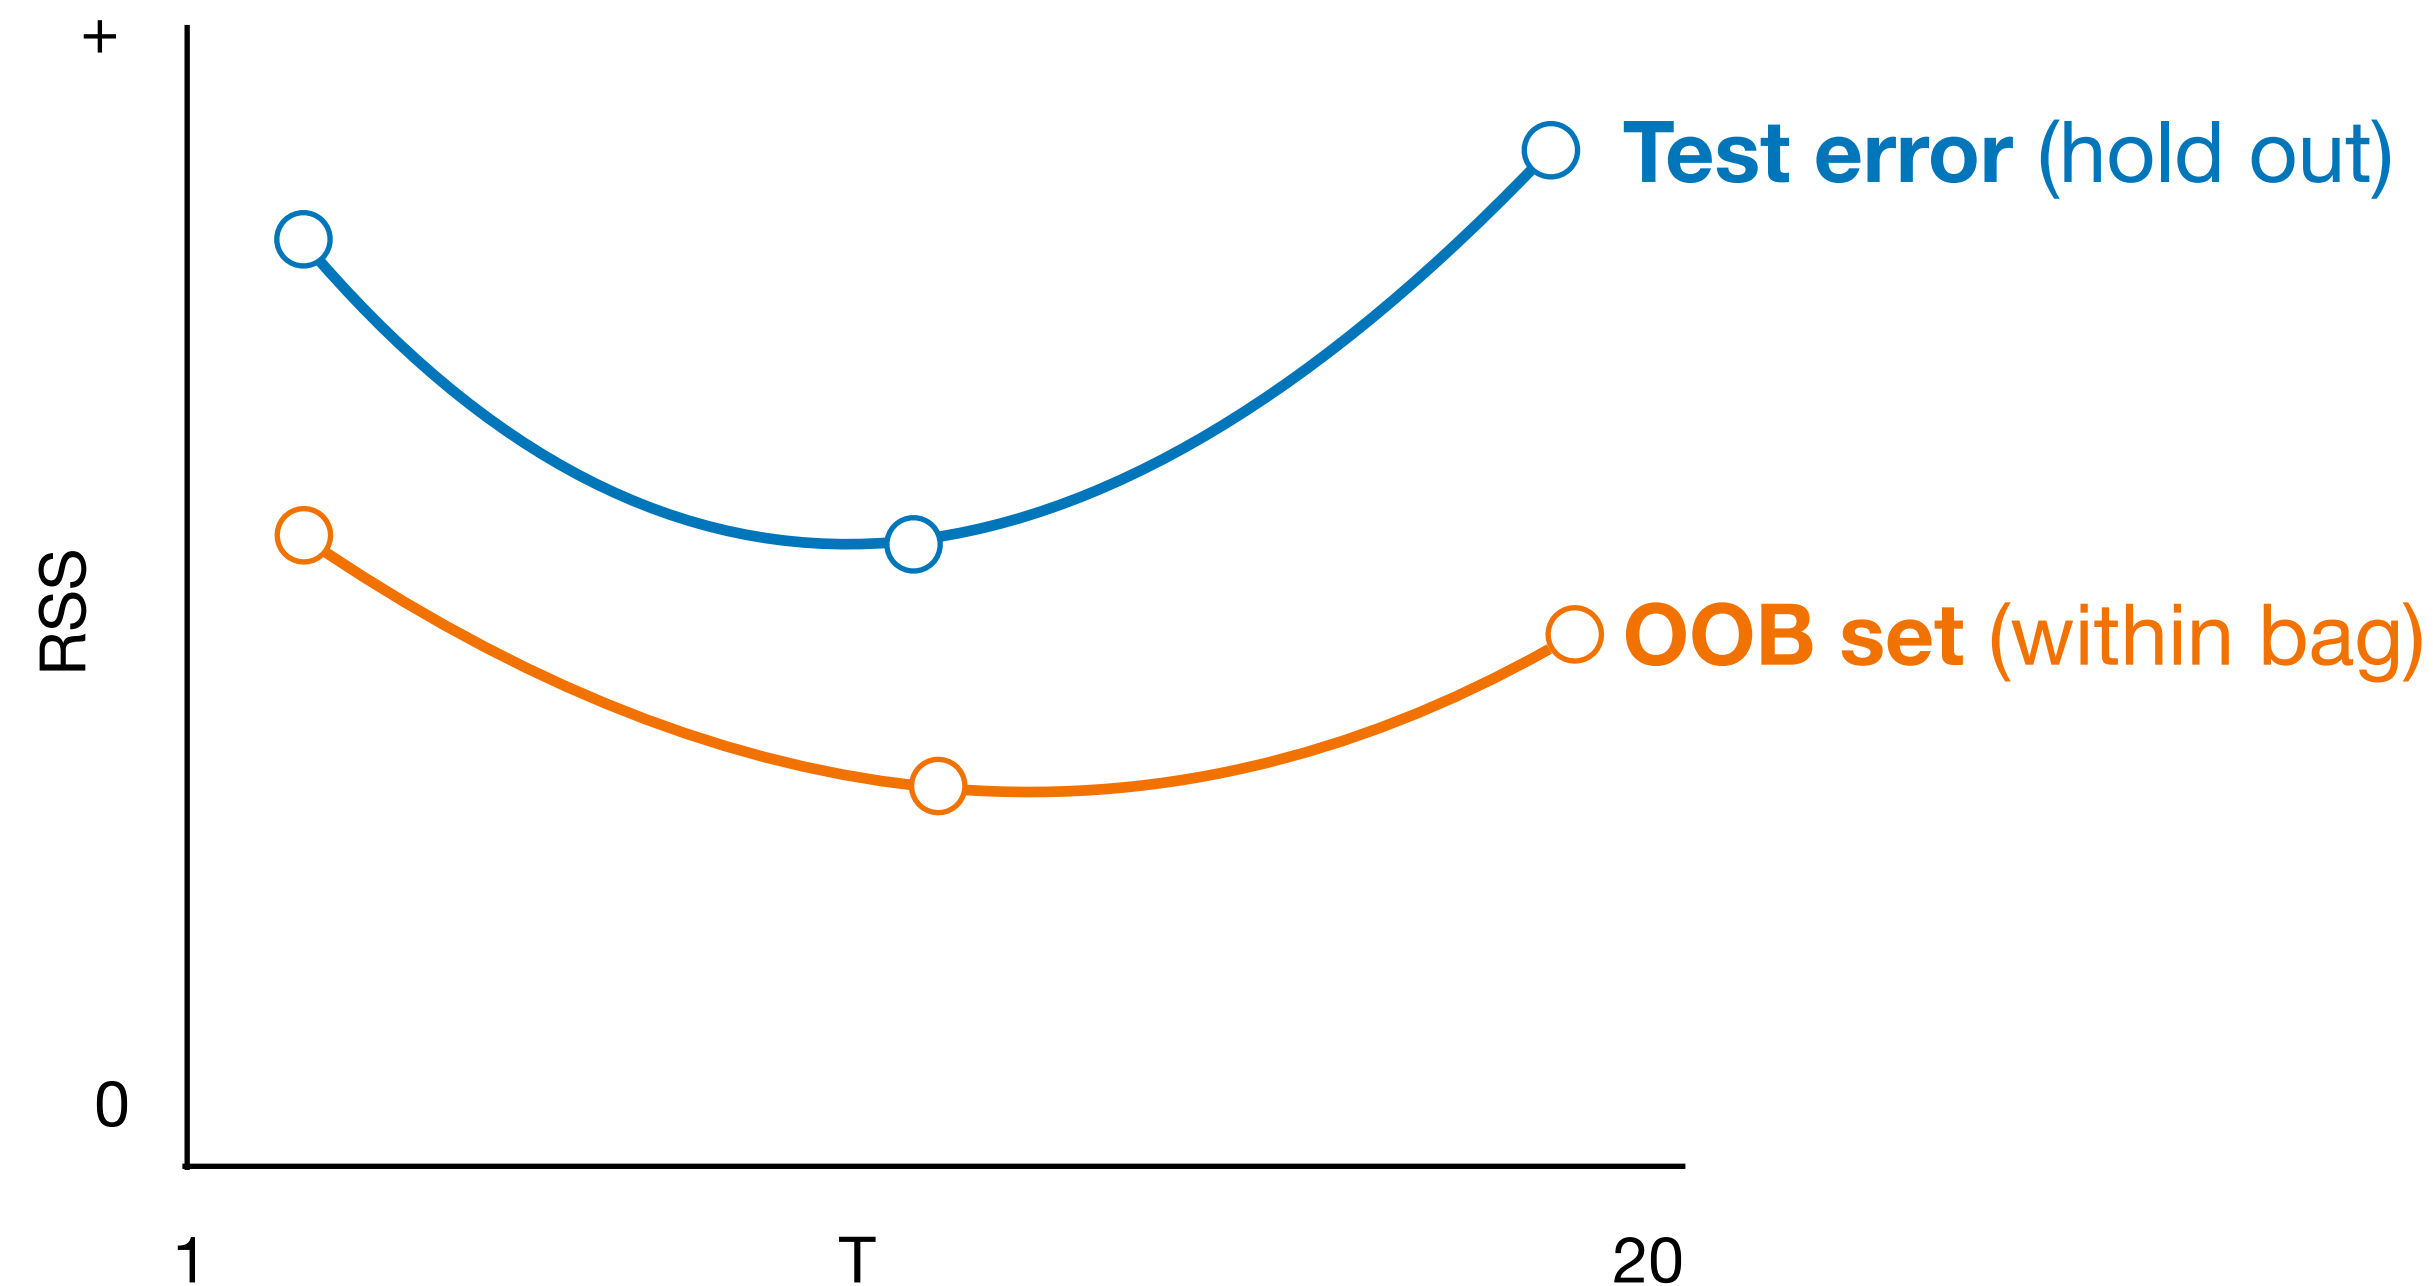
$$\begin{array}{l}
 \overbrace{\begin{pmatrix} y_1 \\ \vdots \\ y_n \end{pmatrix}}^{Y_1^*} = \hat{T}_1\left(\overbrace{\begin{pmatrix} x_{1,1} & \cdots & x_{1,p} \\ \vdots & & \\ x_{n,1} & \cdots & x_{n,p} \end{pmatrix}}^{X_1^*}\right) \rightarrow \hat{f}^1(X_1^*) = \hat{T}_1(X_1^*) \rightarrow \begin{array}{c} \text{---} \\ | \\ \text{---} \end{array} \\
 \vdots \\
 \overbrace{\begin{pmatrix} y_1 \\ \vdots \\ y_n \end{pmatrix}}^{Y_B^*} = \hat{T}_B\left(\overbrace{\begin{pmatrix} x_{1,1} & \cdots & x_{1,p} \\ \vdots & & \\ x_{n,1} & \cdots & x_{n,p} \end{pmatrix}}^{X_B^*}\right) \rightarrow \hat{f}^B(X_B^*) = \hat{T}_B(X_B^*) \rightarrow \begin{array}{c} \text{---} \\ | \\ \text{---} \end{array}
 \end{array}
 \quad \left| \quad \begin{array}{l} \hat{y}_i = \frac{1}{B} \sum_{b=1}^B \hat{f}^b(x_i) \\ \\ = \frac{1}{B} \sum_{b=1}^B \hat{T}_b(x_i) \end{array} \right.$$

Out-of-bag (OOB) error

Each pull of the bootstrap leaves out a few observations (on average 1/3 of observations are left out). Use these as a test set on bagged model.



Out-of-bag (OOB) error




OOB error is usually lower than traditional hold out set error, while maintaining independence between the test and training data.

Variable importance

Goal: On each pull from the bag, b , take each predictor variable, j , and calculate goodness-of-fit using just that variable. Average across all pulls, B .

- Regression = RSS
 - Classification = G
- Variable Importance measures


Approximate estimate of importance.



Random forests


Problem with bagging trees

Because of their flexibility, a few important factors will drive splits in all iterations of the bootstrap, leading to \uparrow correlations across models.




Random forests

Solution: On each split allow only a subset of predictor variables, m , out of all variables, p , to be included in the model.




Tree on iteration b



$$m = \sqrt{p}$$


...



Forest of random trees

$$\begin{array}{l}
 \overbrace{\begin{pmatrix} y_1 \\ \vdots \\ y_n \end{pmatrix}}^{Y_1^*} = \hat{T}_1\left(\overbrace{\begin{pmatrix} x_{1,1} & \cdots & x_{1,p} \\ \vdots & & \\ x_{n,1} & \cdots & x_{n,p} \end{pmatrix}}^{X_1^*}\right) \rightarrow \hat{f}^1(X_1^*) = \hat{T}_1(X_1^*) \rightarrow \begin{array}{c} \text{---} \\ | \quad \text{---} \\ | \quad | \end{array} \\
 \vdots \\
 \overbrace{\begin{pmatrix} y_1 \\ \vdots \\ y_n \end{pmatrix}}^{Y_B^*} = \hat{T}_B\left(\overbrace{\begin{pmatrix} x_{1,1} & \cdots & x_{1,p} \\ \vdots & & \\ x_{n,1} & \cdots & x_{n,p} \end{pmatrix}}^{X_B^*}\right) \rightarrow \hat{f}^B(X_B^*) = \hat{T}_B(X_B^*) \rightarrow \begin{array}{c} \text{---} \\ | \quad \text{---} \\ | \quad | \end{array}
 \end{array}
 \quad \left. \vphantom{\begin{array}{l} \overbrace{\begin{pmatrix} y_1 \\ \vdots \\ y_n \end{pmatrix}}^{Y_1^*} \\ \vdots \\ \overbrace{\begin{pmatrix} y_1 \\ \vdots \\ y_n \end{pmatrix}}^{Y_B^*} \end{array}} \right\} \begin{array}{l}
 \hat{y}_i = \frac{1}{B} \sum_{b=1}^B \hat{f}^b(x_i) \\
 = \frac{1}{B} \sum_{b=1}^B \hat{T}_b(x_i)
 \end{array}$$

Random vs. static forests



- Random forests (RF) forces weaker predictor variables to contribute to predictions.
- On average, a single strong predictor will be in the “approved list” only $\frac{p-m}{p}$ times.
- Improves test accuracy when your predictor variables are highly correlated.

Boosting

Boosting

Bagging & RF: $\hat{f}^1(X_1^*), \dots, \hat{f}^B(X_B^*)$
independent

Boosting: $\hat{f}^1(X_1^*) \rightarrow \hat{f}^2(X_2^*) \rightarrow \dots \rightarrow \hat{f}^B(X_B^*)$
sequential

Let each model inform the next so that you *boost* the overall variance explained by the collective set.

Boosting algorithm

Goal: “Eat up” all the residual variance that you can.


Step 1: Start with a null model, $\hat{f}(X) = 0$, such that $r_i = y_i - \hat{y}_i = y_i$.

Step 2: Run B iterations where, on each iteration b :

- Calculate a new model $\hat{f}^b(X)$ using the objective function.

$$\min \sum_{i=1}^n (r_i - \hat{f}^b(x_i))^2.$$

- Update the previous model with the current model

$$\hat{f}(X) \leftarrow \hat{f}(X) + \lambda \hat{f}^b(X).$$


sparsity parameter

- Update the residuals $r_i \leftarrow r_i - \lambda \hat{f}^b(X)$.

Step 3: Output the boosted model $\hat{f}(X) = \sum_{b=1}^B \lambda \hat{f}^b(X)$

Parameters to tune

Boosting relies on 3 free parameters that have to be tuned.

1. $B \rightarrow$ number of models generated
2. $\lambda \rightarrow$ sparsity constraint
3. $d \rightarrow$ number of splits (if using trees)

Care must be taken with cross validation sets when selecting these parameters.

Take home message

- Bagging, random forests, and boosting are very powerful methods that improve overall prediction accuracy of high variance methods (e.g., decision trees), but at the expense of interpretability.